

ASDM Workshop Week 4 : Clustering with R

- Dr. Mo Saraee, Charith Silva

Clustering analysis is widely used in many fields. Traditionally clustering is regarded as unsupervised learning for its lack of a class label or a quantitative response variable, which in contrast is present in supervised learning such as classification and regression.

Why Clustering analysis is unsupervised? Because, the target variable is not present. The model is trained based on given input variables which attempt to discover intrinsic groups (or clusters).

Clustering is a division of data into groups of similar objects. Representing the data by fewer clusters necessarily loses certain fine details, but achieves simplification. It models data by its clusters. Data modeling puts clustering in a historical perspective rooted in mathematics, statistics, and numerical analysis. From a machine learning perspective clusters correspond to hidden patterns, the search for clusters is unsupervised learning, and the resulting system represents a data concept. From a practical perspective clustering plays an outstanding role in data mining applications such as scientific data exploration, information retrieval and text mining, spatial database applications, Web analysis, CRM, marketing, medical diagnostics, computational biology (Berkhin,2006).

Some other applications of clustering are:

- Marketing: Help marketers discover distinct groups in their customer bases, and then use this knowledge to develop targeted marketing programs.
- Land use: Identification of areas of similar land use in an earth observation database.
- Insurance: Identifying groups of motor insurance policy holders with a high average claim cost.
- City-planning: Identifying groups of houses according to their house type, value, and geographical location.

Part 1: Exercise

We will be using **costpercompany.csv** dataset in this workshop which can be downloaded from Blackboard. The dataset contains several different financial indicators about the given companies.

Data Explanation

Variable	Explanation
Company	Name of the company
surcharges	Surcharge cost per day
RoR	Return on Risk
dailycost	Daily miscellaneous cost
costwithload	Cost related to loading
costofDemand	Demand cost
Sales	Sales of company
WearandTear	Natural wear and tear cost
Fcost	Fuel cost

Implementation:

In this workshop, we will be using "cluster" packages which use K-means clustering Algorithm. K-means (MacQueen, 1967) is one of the simplest unsupervised learning algorithms that solve the well-known clustering problem. The procedure follows a simple and easy way to classify a given data set through a certain number of clusters (assume k clusters). We will also be using hierarchical clustering for basic understanding.

1. Download **costpercompany.csv** dataset from Blackboard and save it to a folder on your F: drive (eg: F:\ASDM\Week4). Open it using excel to get a rough idea about the dataset, e.g, attributes and their values.
2. Start RStudio.
3. Change the working directory
File → More → Go To Working Directory...

In the Go To Working Directory dialogue, navigate and select the folder where you saved your data file eg: F:\ASDM\Week4. Click OK.

Click Set as Working Directory option

or

using R commands as follow:

```
mypath = "F:\\ASDM\\Week4" # you need to change the string to your
                             directory
setwd(mypath)               # set working directory
getwd()                     # check if the working directory has
                             changed correctly
```

4. Open a new R script window:
File → New File → R script

5. Read the data file

```
cost_data <- read.csv("costpercompany.csv", header= TRUE)
```

The data is read into the data frame named "cost_data"

6. Inspect the dataset in R

Once the file has been imported to R, we often want to do few things to explore the dataset:

```
names(cost_data)
head(cost_data)
tail(cost_data)
summary(cost_data)
str(cost_data)
```

```
> tail(cost_data)
  Company surcharges  RoR dailycost costwithload costofDemand Sales WearandTear Fcost
24 Commonwealth  1.02 11.20    168      56.0         0.3  6423      34.3 0.700
25 Central        1.43 15.40    113      53.0         3.4  9212       0.0 1.058
26 CA Gas         1.95  1.86     49     -39.4         6.4 15280      12.3 1.140
27 BP            3.90 21.16    370      72.0        16.4 40008      53.4 2.610
28 Boston        0.89 10.30    202      57.9         2.2  5088      25.3 1.555
29 Arizona        1.06  9.20    151      54.4         1.6  9077       0.0 0.628
> summary(cost_data)
  Company      surcharges      RoR      dailycost      costwithload      costofDemand      Sales      WearandTear      Fcost
Arizona : 1   Min.   :0.750   Min.   : 1.86   Min.   : 49.0   Min.   : -49.80   Min.   : -2.200   Min.   : 3300   Min.   : 0.00   Min.   : -0.012
Boston  : 1   1st Qu.:1.050   1st Qu.: 9.20   1st Qu.:148.0   1st Qu.: 51.50   1st Qu.:  2.200   1st Qu.: 6650   1st Qu.: 0.00   1st Qu.:  0.636
BP      : 1   Median :1.150   Median :10.58   Median :173.0   Median : 56.00   Median :  3.500   Median : 9673   Median : 8.30   Median :  1.108
CA Gas  : 1   Mean   :1.403   Mean   :10.49   Mean  :172.4   Mean   : 46.86   Mean   :  4.907   Mean  :13025   Mean   :16.11   Mean   :  1.174
Central : 1   3rd Qu.:1.430   3rd Qu.:12.20   3rd Qu.:199.0   3rd Qu.: 60.00   3rd Qu.:  7.200   3rd Qu.:15651   3rd Qu.:26.70   3rd Qu.:  1.652
Commonwealth: 1   Max.   :3.900   Max.   :21.16   Max.   :370.0   Max.   : 72.00   Max.   :16.400   Max.   :40008   Max.   :53.40   Max.   :  2.610
(Other) :23
> str(cost_data)
'data.frame':   29 obs. of  9 variables:
 $ Company      : Factor w/ 29 levels "Arizona ","Boston ",...: 29 28 27 26 25 24 23 22 21 20 ...
 $ surcharges   : num  2.7 1.2 1.07 1.04 1.16 1.05 1.95 0.76 1.16 0.96 ...
 $ RoR          : num  9.36 11.8 9.3 8.6 11.7 ...
 $ dailycost    : int  222 148 174 204 104 150 185 136 252 164 ...
 $ costwithload : num  12.1 59.9 54.3 61 54 56.7 36 61.9 56 62.2 ...
 $ costofDemand : num  12.9 3.5 5.9 3.5 -2.1 2.7 8.2 9 9.2 -0.1 ...
 $ Sales        : int  32721 7287 10093 6650 13507 10140 20004 5714 15991 6468 ...
 $ WearandTear  : num  12.3 41.1 26.6 0 0 0 26.7 8.3 0 0.9 ...
 $ Fcost        : num  1.908 0.702 1.306 2.116 0.636 ...
```

7. Check the dimension and number of points of the "cost_data" dataset

```
nrow(cost_data)
ncol(cost_data)
dim(cost_data)
```

```
> nrow(cost_data)
[1] 29
> ncol(cost_data)
[1] 9
> dim(cost_data)
[1] 29 9
```

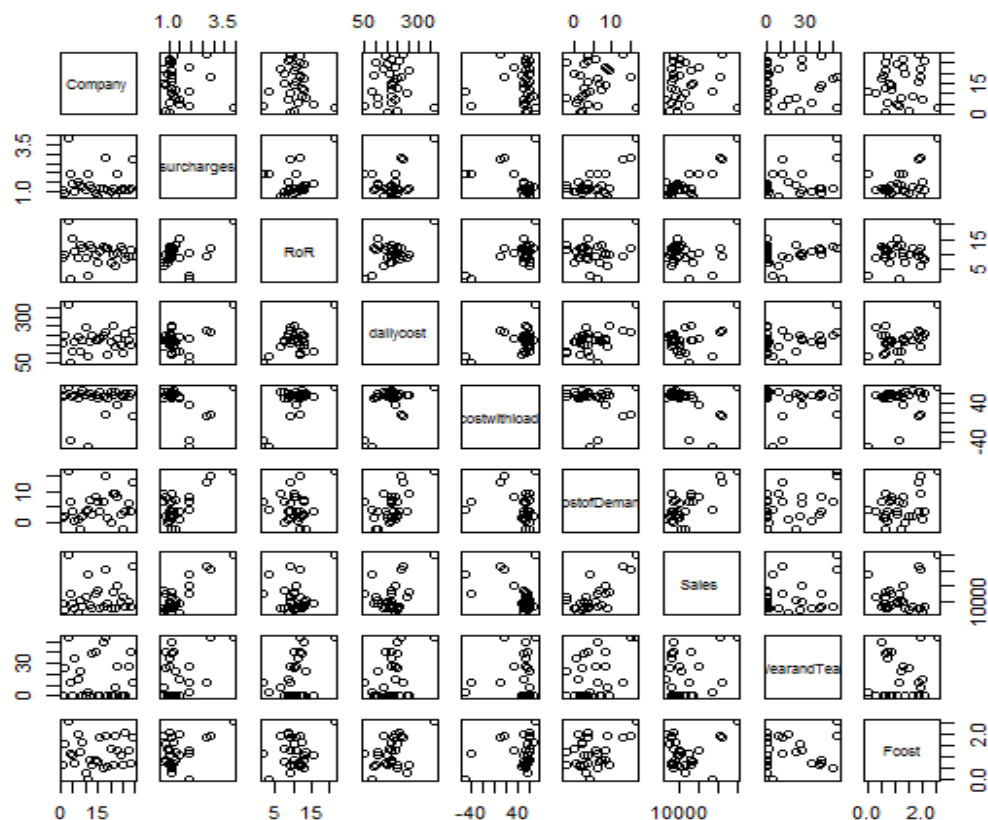
8. Install and activate "cluster" package.

#cluster package is a powerful tool for cluster analysis

```
install.packages("cluster") # install "cluster" package
library(cluster)            # activate "cluster" package
```


9. Following command can be used to create scatterplot matrix to compare the variables

```
pairs(cost_data)
```



10. Use the following line of code to plot and understand the relationship between the Sales and Fuel cost of the company:


```
plot(Fcost~ Sales, data = cost_data)
```



11. The result does not tell you anything about the companies related to each cluster. Hence, using the following line of code to see the names of the companies:


```
with(cost_data,text(Fcost ~ Sales, labels= Company))
```

Result:



12. Remove the overlap of names of the companies:

```
plot(Fcost~ Sales, data = cost_data)
with(cost_data,text(Fcost ~ Sales, labels= Company,pos=4,cex=.6))
```



13. Normalization of the dataset is among the preprocessing processes in data exploration, in which the attribute data are scaled to fall in a small specified range.

In this dataset some variables are in decimal points (Fuel-cost) and some are with very high values like thousands (Sales), so these values will be dominating in the analysis.

Normalization before clustering is specifically needed for distance matrix, like the Euclidian distance that are sensitive to variations within the magnitude or scales from the attributes.

#normalise function

```
normalise <- function(df)
{
  return(((df- min(df)) / (max(df)-min(df)) * (1-0)) +0)
}
```

```
head(cost_data)
```

```
> head(cost_data)
  Company surcharges RoR dailycost costwithload costofDemand Sales WearandTear Fcost
1  Yashida      2.70  9.36      222         12.1         12.9 32721      12.3 1.908
2 Wisconsin     1.20 11.80      148         59.9          3.5  7287      41.1 0.702
3  Virginia     1.07  9.30      174         54.3          5.9 10093      26.6 1.306
4   United     1.04  8.60      204         61.0          3.5  6650       0.0 2.116
5   Texas     1.16 11.70      104         54.0         -2.1 13507       0.0 0.636
6 Southern     1.05 12.60      150         56.7          2.7 10140       0.0 1.108
```

```
#Normalise the dataset using above normalise function
```

```
company<-cost_data[,1]
```

```
cost_data_n<-cost_data[,2:9] #remove company column before normalise
```

```
cost_data_n<-as.data.frame(lapply(cost_data_n,normalise))
```

```
cost_data_n$Company<-company #add company column after normalise
```

```
#rearrange the columns in the dataset after normalising
```

```
cost_data_n<-cost_data_n[,c(9,1,2,3,4,5,6,7,8)]
```

```
head(cost_data_n)
```

```
> head(cost_data_n)
  Company surcharges RoR dailycost costwithload costofDemand Sales WearandTear Fcost
1  Yashida 0.61904762 0.3886010 0.5389408  0.5082102  0.811827957 0.80148741  0.2303371 0.7322654
2 Wisconsin 0.14285714 0.5150259 0.3084112  0.9006568  0.306451613 0.10861393  0.7696629 0.2723112
3  Virginia 0.10158730 0.3854922 0.3894081  0.8546798  0.435483871 0.18505503  0.4981273 0.5026697
4   United 0.09206349 0.3492228 0.4828660  0.9096880  0.306451613 0.09126076  0.0000000 0.8115942
5   Texas 0.13015873 0.5098446 0.1713396  0.8522167  0.005376344 0.27805928  0.0000000 0.2471396
6 Southern 0.09523810 0.5564767 0.3146417  0.8743842  0.263440860 0.18633540  0.0000000 0.4271548
```

14. Since we have normalized our data, now we can calculate the distance matrix using Euclidean Distance. The choice of distance measures is very important, as it has a strong influence on the clustering results. For most common clustering the default distance measure is the Euclidean distance.

```
#Dist() function computes and returns the distance matrix computed by using the specified distance measure to compute the distances between the rows of a data matrix. Dist() function accepts only numeric data as an input.
```

```
# Note that, allowed values for the option method include one of:
"euclidean", "maximum", "manhattan", "canberra", "binary" and
"minkowski".
```

```
#choose distance method and create distance matrix
```

```
distance <- dist(cost_data_n,method = "euclidean",)
```

```
# In this matrix, the value represents the distance between
companies.
```

```
print(distance)
```

```
> distance
      1      2      3      4      5      6      7      8      9     10     11     12     13     14
2  1.3779637
3  1.0900706 0.4455350
4  1.1918947 1.0305857 0.6557062
5  1.3910132 0.9085196 0.8014188 0.8020101
6  1.2150186 0.8411797 0.5987688 0.5099920 0.3851189
7  0.6889819 0.6889819 0.4676341 0.8755935 0.9515602 0.7856404
8  1.1775998 0.9378339 0.5454479 0.4545513 0.9312438 0.6405588 0.8114437
9  1.0173962 0.9832388 0.6994714 0.7606013 0.8183244 0.5847719 0.7254177 0.7724816
10 1.3409603 0.9071903 0.6371691 0.3843243 0.4999180 0.3631547 0.9142693 0.5941891 0.7515437
11 1.3470049 0.8522907 0.7207823 0.7590054 0.2352620 0.2943494 0.8920285 0.8087363 0.7168157 0.4966276
12 0.8406961 1.3059485 1.2161691 1.6068993 1.7801013 1.6184633 0.8552425 1.5004362 1.4742499 1.7338989 1.7306672
13 1.4156067 0.3147726 0.6078059 1.2212377 1.1688037 1.0627062 0.7594840 1.1043016 1.0680394 1.1501769 1.0955256 1.2163865
14 1.1919754 0.9562812 0.6146599 0.1803758 0.6938388 0.3721955 0.8324323 0.4633696 0.7152954 0.3550891 0.6296577 1.5944125 1.1600842
15 1.0890280 0.9345940 0.6283006 0.6740632 0.6173063 0.4549315 0.7450020 0.6374166 0.3676625 0.5678501 0.5286573 1.5516763 1.0950898 0.6327377
```

```
#Round the distance figures to 3 decimals.
```

```
print(distance,digits=3)
```


```
> print(distance,digits=3)
      1      2      3      4      5      6      7      8      9     10     11     12     13     14     15     16     17     18     19     20     21     22     23     24     25     26     27
2  1.378
3  1.090 0.446
4  1.192 1.031 0.656
5  1.391 0.909 0.801 0.802
6  1.215 0.841 0.599 0.510 0.385
7  0.689 0.689 0.468 0.876 0.952 0.786
8  1.178 0.938 0.545 0.455 0.931 0.641 0.811
9  1.017 0.983 0.699 0.761 0.818 0.585 0.725 0.772
10 1.341 0.907 0.637 0.384 0.500 0.363 0.914 0.594 0.752
11 1.347 0.852 0.721 0.759 0.235 0.294 0.892 0.809 0.717 0.497
12 0.841 1.306 1.216 1.607 1.780 1.618 0.855 1.500 1.474 1.734 1.731
13 1.416 0.315 0.608 1.221 1.169 1.063 0.759 1.104 1.068 1.150 1.096 1.216
14 1.192 0.956 0.615 0.180 0.694 0.372 0.832 0.463 0.715 0.355 0.630 1.594 1.160
15 1.089 0.935 0.628 0.674 0.617 0.455 0.745 0.637 0.368 0.568 0.529 1.552 1.095 0.633
16 1.174 0.406 0.413 1.049 1.049 0.914 0.545 0.869 0.871 1.008 0.972 1.089 0.400 1.003 0.841
17 1.384 0.192 0.467 1.011 0.905 0.826 0.708 0.975 0.934 0.891 0.856 1.340 0.308 0.944 0.923 0.465
18 1.139 0.853 0.610 0.622 0.607 0.305 0.733 0.639 0.426 0.574 0.445 1.539 1.017 0.489 0.456 0.874 0.839
19 1.266 1.485 1.361 1.538 1.210 1.315 1.195 1.499 1.276 1.388 1.191 1.699 1.596 1.507 1.135 1.434 1.481 1.340
20 1.244 0.919 0.726 0.768 0.591 0.497 0.850 0.878 0.372 0.594 0.552 1.664 1.060 0.716 0.406 0.934 0.847 0.483 1.235
21 1.224 0.913 0.624 0.311 0.573 0.272 0.831 0.581 0.695 0.323 0.534 1.619 1.132 0.166 0.616 0.993 0.897 0.438 1.493 0.644
22 1.349 0.572 0.572 0.789 0.529 0.556 0.774 0.884 0.983 0.610 0.580 1.506 0.856 0.691 0.842 0.800 0.601 0.744 1.423 0.838 0.596
23 1.218 0.805 0.528 0.398 0.813 0.612 0.769 0.598 0.915 0.468 0.792 1.467 1.008 0.411 0.836 0.923 0.789 0.735 1.505 0.858 0.456 0.597
24 1.425 0.250 0.457 0.928 0.757 0.729 0.757 0.521 0.940 0.751 0.735 1.454 0.496 0.861 0.862 0.558 0.211 0.806 1.437 0.796 0.802 0.452 0.703
25 1.218 0.867 0.691 0.644 0.440 0.242 0.802 0.754 0.700 0.556 0.318 1.600 1.088 0.497 0.618 0.972 0.872 0.331 1.329 0.651 0.398 0.571 0.709 0.793
26 1.118 1.287 1.095 1.250 1.182 1.174 0.989 1.135 1.239 1.184 1.113 1.475 1.419 1.237 1.068 1.246 1.319 1.200 0.637 1.256 1.278 1.242 1.170 1.271 1.192
27 1.388 1.932 1.856 2.085 2.320 2.130 1.534 2.103 1.957 2.271 2.302 0.999 1.832 2.071 2.137 1.768 1.933 2.033 2.486 2.159 2.069 2.045 2.003 2.057 2.098 2.341
28 1.278 0.526 0.304 0.569 0.794 0.597 0.694 0.619 0.831 0.545 0.752 1.407 0.711 0.543 0.761 0.641 0.492 0.693 1.518 0.770 0.546 0.505 0.368 0.427 0.715 1.247 1.968
29 1.306 0.841 0.649 0.643 0.324 0.279 0.855 0.727 0.585 0.358 0.243 1.710 1.066 0.554 0.406 0.933 0.819 0.412 1.197 0.370 0.473 0.629 0.703 0.701 0.434 1.115 2.274
```

15. Visualise the distance matrices using a function called **fviz_dist()** in **factoextra** package

```
install.packages("factoextra") # install "factoextra" package
library(factoextra)           # activate "factoextra" package
```



```
fviz_dist(distance)
```



The colour level is proportional to the value of the dissimilarity between observations:

Red: high similarity (ie: low dissimilarity)

Blue: low similarity

16. Add company name to the distance matrix

```
#inspect the first few observations
```

```
head(cost_data_n)
```

```
> head(cost_data_n)
  Company surcharges      RoR  dailycost costwithload costofDemand      Sales WearandTear      Fcost
1  Yashida 0.61904762 0.3886010 0.5389408  0.5082102 0.811827957 0.80148741 0.2303371 0.7322654
2 Wisconsin 0.14285714 0.5150259 0.3084112  0.9006568 0.306451613 0.10861393 0.7696629 0.2723112
3  Virginia 0.10158730 0.3854922 0.3894081  0.8546798 0.435483871 0.18505503 0.4981273 0.5026697
4   United 0.09206349 0.3492228 0.4828660  0.9096880 0.306451613 0.09126076 0.0000000 0.8115942
5   Texas 0.13015873 0.5098446 0.1713396  0.8522167 0.005376344 0.27805928 0.0000000 0.2471396
6 Southern 0.09523810 0.5564767 0.3146417  0.8743842 0.263440860 0.18633540 0.0000000 0.4271548
```

```
#rownames() function retrieve the row names of a matrix-like object.
```

```
#Set company names as row names
```

```
rownames(cost_data_n)<-cost_data_n$Company
```

```
#remove Company column from the dataset
```

```
cost_data_n$Company<-NULL
```


```
#inspect the top observations
```

```
head(cost_data_n)
```

```
> head(cost_data_n)
  surcharges      RoR  dailycost costwithload costofDemand      Sales  WearandTear      Fcost
Yashida  0.61904762 0.3886010 0.5389408   0.5082102  0.811827957 0.80148741  0.2303371 0.7322654
Wisconsin 0.14285714 0.5150259 0.3084112   0.9006568  0.306451613 0.10861393  0.7696629 0.2723112
Virginia 0.10158730 0.3854922 0.3894081   0.8546798  0.435483871 0.18505503  0.4981273 0.5026697
United   0.09206349 0.3492228 0.4828660   0.9096880  0.306451613 0.09126076  0.0000000 0.8115942
Texas    0.13015873 0.5098446 0.1713396   0.8522167  0.005376344 0.27805928  0.0000000 0.2471396
Southern 0.09523810 0.5564767 0.3146417   0.8743842  0.263440860 0.18633540  0.0000000 0.4271548
```


```
distance <- dist(cost_data_n,method = "euclidean")
```

```
fviz_dist(distance)
```



17. Hierarchical clustering

In simple words, hierarchical clustering tries to create a sequence of nested clusters to explore deeper insights from the data. Most commonly used hierarchical clustering method is **Agglomerative Clustering**. It starts with treating every observation as a cluster. Then, it merges the most similar observations into a new cluster. This process continues until all the observations are merged into one cluster. It uses a bottoms-up approach. This hierarchy of clusters is graphically presented using a Dendrogram (shown below).



#choose previously created intercluster distance matrix and perform Hierarchical clustering using hclust() function


```
cost_data.hclust <- hclust(distance)
cost_data.hclust
```

```
> cost_data.hclust <- hclust(distance)
> cost_data.hclust
```


```
Call:
hclust(d = distance)
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 29
```

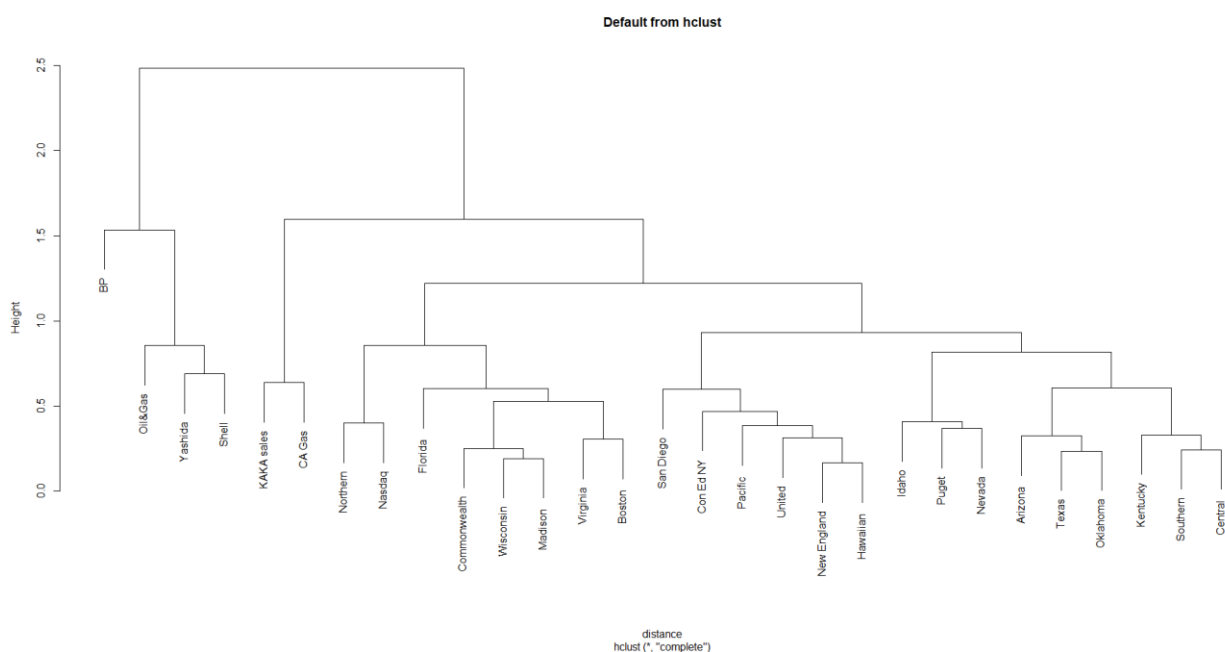
```
plot(cost_data.hclust) # plot the results
```



```
plot(cost_data.hclust, hang=-1)
```



```
plot(cost_data.hclust, labels=cost_data$Company)
```

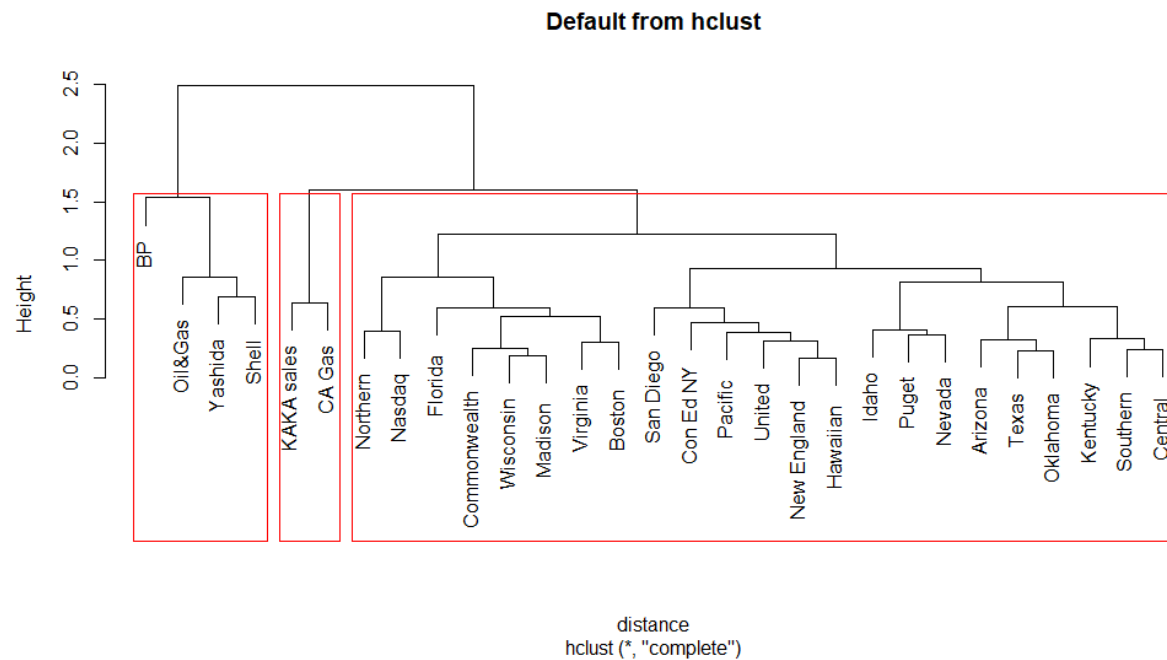


18. plot clusters for chosen number of clusters

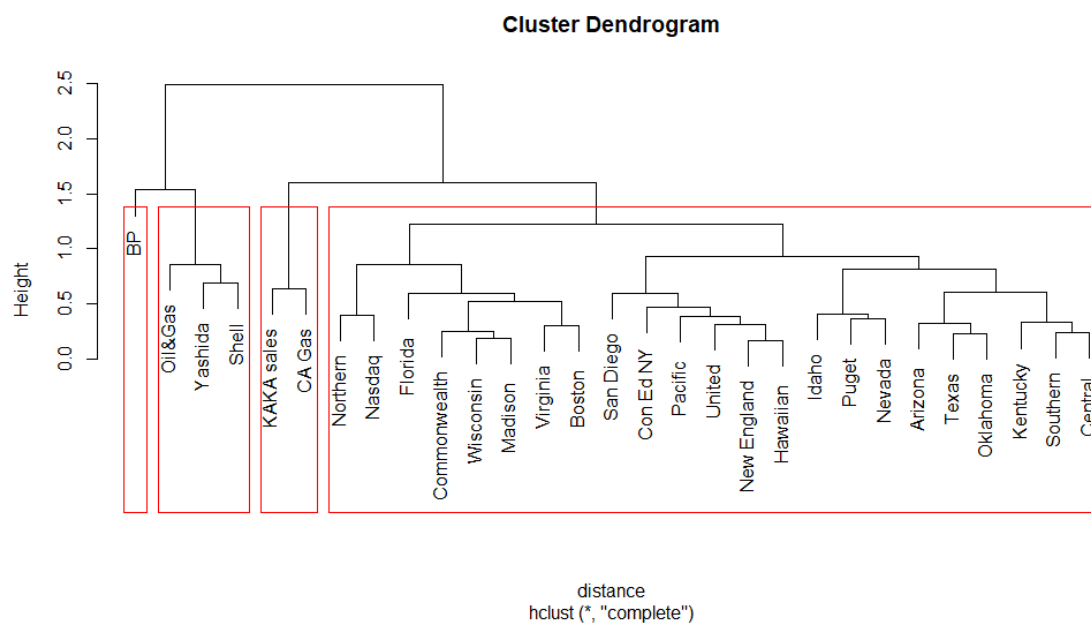
`rect.hclust()` function draws rectangles around the branches of a dendrogram highlighting the corresponding clusters. First the

dendrogram is cut at a certain level, then a rectangle is drawn around selected branches.

```
# draw 3 clusters
rect.hclust(cost_data.hclust, 3)
```



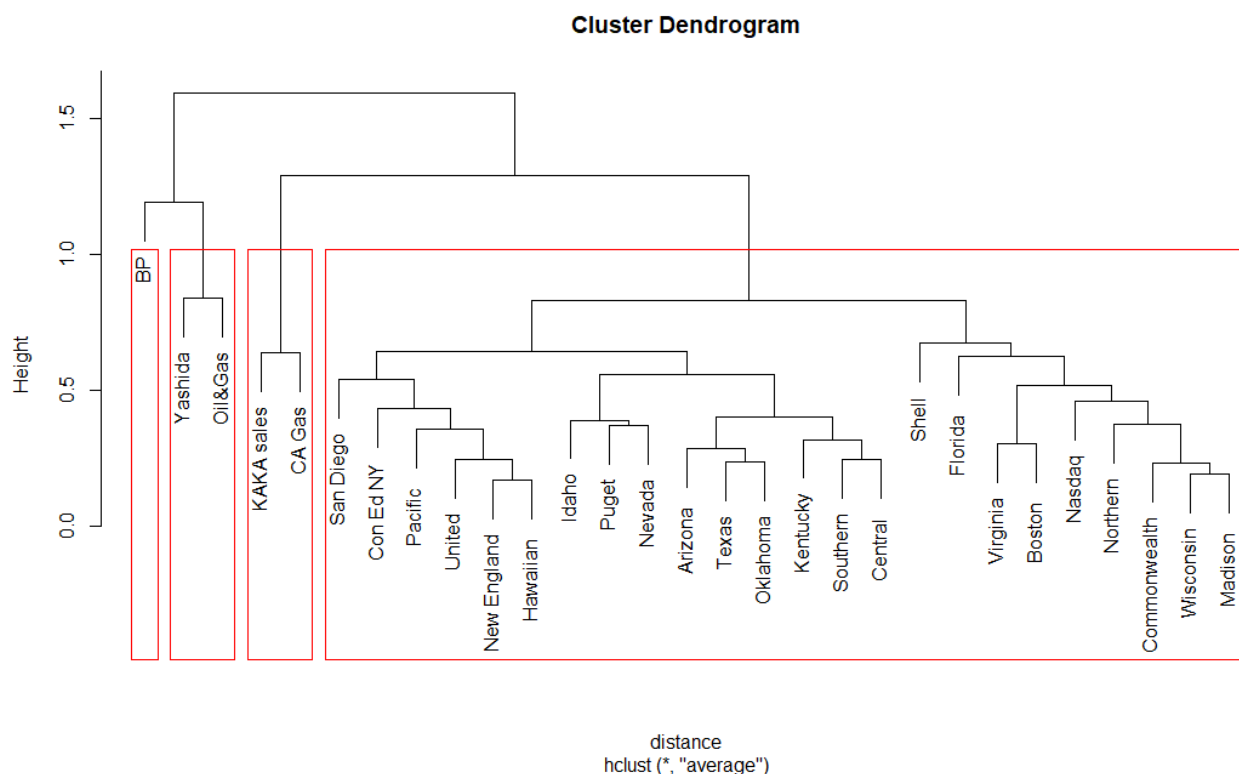
```
# draw 4 clusters
plot(cost_data.hclust, labels=cost_data$Company)
rect.hclust(cost_data.hclust, 4)
```




19. There are a few ways to determine how close two clusters are:

- **Complete linkage clustering:** Find the maximum possible distance between points belonging to two different clusters.
- **Single linkage clustering:** Find the minimum possible distance between points belonging to two different clusters.
- **Average linkage clustering:** Find all possible pairwise distances for points belonging to two different clusters and then calculate the average.
- **Centroid linkage clustering:** Find the centroid of each cluster and calculate the distance between centroids of two clusters.


```
# Cluster using average linkage
hclust.average <- hclust(distance, method = "average")
plot(hclust.average, labels=cost_data$Company)
rect.hclust(hclust.average, 4)
```



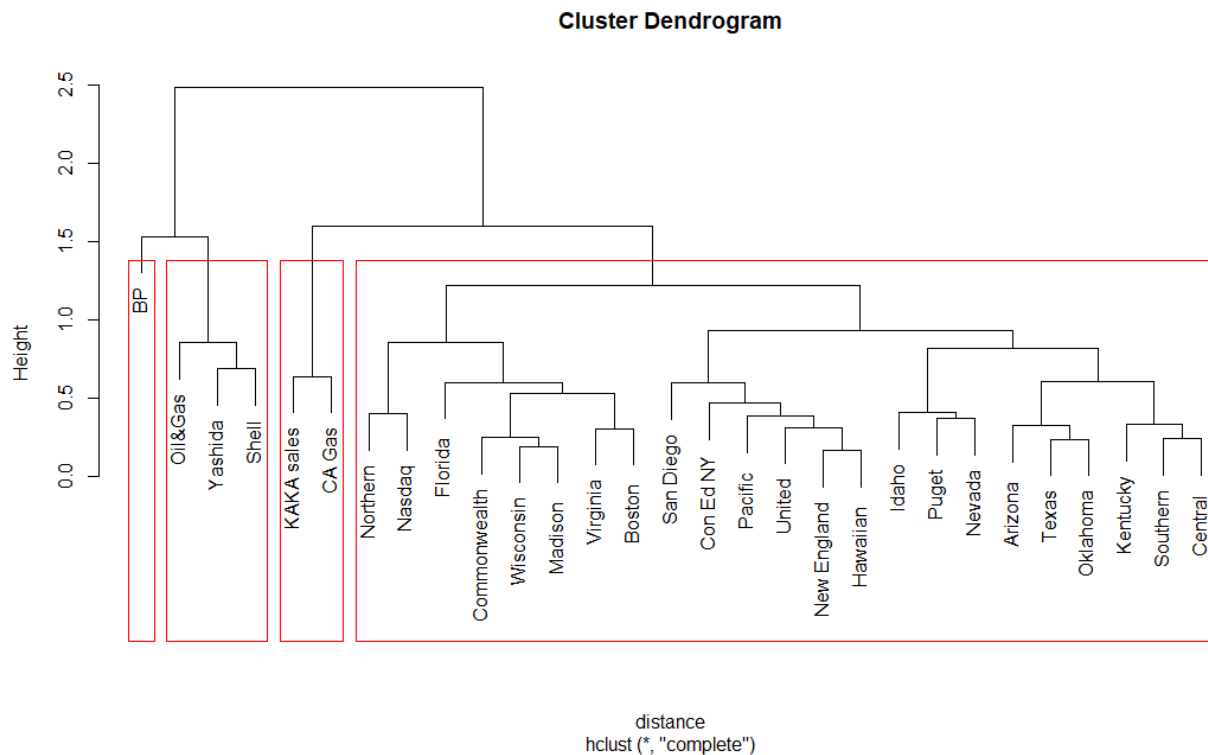
```
# Cluster using single linkage
hclust.single <- hclust(distance, method = "single")
plot(hclust.single, labels=cost_data$Company)
rect.hclust(hclust.single, 4)
```



```
# Cluster using centroid linkage
hclust.centroid<- hclust(distance, method = "centroid")
plot(hclust.centroid, labels=cost_data$Company)
rect.hclust(hclust.centroid, 4)
```



```
# Cluster using complete linkage
hclust.complete <- hclust(distance, method = "complete")
plot(hclust.complete, labels=cost_data$Company)
rect.hclust(hclust.complete, 4)
```



20. Compare the cluster membership

#cutree function() cuts a dendrogram tree into several groups by specifying the desired number of clusters.

```
member.centroid <- cutree(hclust.centroid,4)
member.centroid
member.complete <- cutree(hclust.complete,4)
member.complete

table(member.centroid,member.complete)
```

```
> table(member.centroid,member.complete)
      member.complete
member.centroid  1  2  3  4
1      1  0  0  0
2      1 23  2  0
3      1  0  0  0
4      0  0  0  1
```


21. K-Means clustering

K-means clustering is a type of unsupervised learning technique, which is used when you have unlabeled data (i.e., data without defined categories or groups). The goal of this algorithm is to find groups in the data, with the number of groups represented by the variable K . The algorithm works iteratively to assign each data point to one of K groups based on the features that are provided. Data points are clustered based on feature similarity.

K-Means clustering algorithm can be used for any type of grouping. Some examples of use cases are: (source : Oracle data science)

- Behavioural segmentation:
 - Segment by purchase history
 - Segment by activities on application, website, or platform
 - Define personas based on interests
 - Create profiles based on activity monitoring
- Inventory categorization:
 - Group inventory by sales activity
 - Group inventory by manufacturing metrics
- Sorting sensor measurements:
 - Detect activity types in motion sensors
 - Group images
 - Separate audio
 - Identify groups in health monitoring
- Detecting bots or anomalies:
 - Separate valid activity groups from bots
 - Group valid activity to clean up outlier detection

Use the following line of code to create k means clustering:

```
#k-means clustering is an algorithm used to find homogeneous subgroups in a population
```

#Defining the optimal number of cluster as k-means clustering requires to specify the number of clusters to generate.

#kmeans function perform k-means clustering on a data matrix.

```
kc<-kmeans(cost_data[,-1],3) #k=3
```

```
kc
```

```
> kc<-kmeans(cost_data[,-1],3)
> kc
K-means clustering with 3 clusters of sizes 4, 7, 18

Cluster means:
  surcharges      RoR  dailycost costwithload costofDemand      Sales  WearandTear      Fcost
1    2.845 11.360000  224.2500    12.97500    12.000000 32666.750    30.77500  1.6220000
2    1.280  8.691429  168.2857    38.87143     5.700000 15850.857    11.44286  0.8028571
3    1.130 10.994444  162.5556    57.50000     3.022222  7560.444    14.66667  1.2181667

Clustering vector:
[1] 1 3 3 3 2 3 2 3 2 3 3 1 3 3 2 2 3 3 1 2 3 3 3 3 3 2 1 3 3


Within cluster sum of squares by cluster:
[1] 88992681 33363677 70434368
(between_SS / total_SS =  91.7 %)

Available components:

[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss" "betweenss"
[7] "size"         "iter"         "ifault"
```

22. The result can be plotted using the `clusplot()` function

```
clusplot(cost_data, kc$cluster, color=TRUE, shade=TRUE, lines=0)
```



Part 2: Exercise

In this exercise we will practice how to implement K-Means clustering in real-world dataset.

The dataset, **crimes-2017-18.csv** can be downloaded from Blackboard. Dataset contains quarterly statistics on various crime figures in England and Wales.

Source : <https://www.gov.uk/government/statistics/police-recorded-crime-open-data-tables>

Dataset Explanation

FINANCIAL YEAR	
Possible values	Various
Combined with the Financial Quarter column, this identifies the period during which offences took place. Each financial year runs from April to March.	
FINANCIAL QUARTER	
Possible values	Various

Combined with Financial Year column, this identifies the period during which offences took place. Quarter 1 runs from April-June, Quarter 2 from July-September, Quarter 3 from October-December, and Quarter 4 from January-March.	
FORCE NAME	
Possible values	Various
This column identifies the police force area in which offences took place. The reference table 'PRC Geog reference table.csv' shows how these areas map up to regions within England and Wales.	
CSP NAME (CSP tables only)	
Possible values	Various
This column identifies the Community Safety Partnership in which offences took place. This is a geographic area within a Police Force, and generally corresponds to Local Authority boundaries. The reference table 'PRC Geog reference table.csv' shows how these areas map up to Police Forces and regions within England and Wales.	
OFFENCE DESCRIPTION	
Possible values	Various
This column provides a description of the offence covered by each Offence Code value.	
OFFENCE GROUP	
Possible values	Various
This column identifies the offence group within which the Offence Code falls. Each groups also consists of Offence Sub-groups , which in turn consist of Offence Codes .	

OFFENCE SUBGROUP	
Possible values	Various
This column identifies the offence sub-group within which the Offence Code falls. These sub-groups contain Offence Codes .	
OFFENCE CODE	
Possible values	Various
This column identifies the specific offence code used by the police and the Home Office to classify offences. The reference table 'Ref-Offence.csv' shows descriptions of these codes, as well as the offence groups that they map up to.	
NUMBER OF OFFENCES	
Possible values	Various
This column contains the total number of police recorded crimes for the specified Offence Code , CSP Name/Force Name and time period (Financial Year and Financial Quarter).	

1. Read the data file

```
crimes <- read.csv("crimes-2017-18.csv", header= TRUE)
```

2. Inspect the dataset

```
names(crimes)
head(crimes)
tail(crimes)
summary(crimes)
str(crimes)
```

```
nrow(crimes)
ncol(crimes)
dim(crimes)
```

```
> str(crimes)
'data.frame': 191260 obs. of 9 variables:
 $ Financial.Year : Factor w/ 1 level "2017/18": 1 1 1 1 1 1 1 1 1 1 ...
 $ Financial.Quarter : int 1 1 1 1 1 1 1 1 1 1 ...
 $ Force.Name : Factor w/ 43 levels "Avon & Somerset",...: 1 1 1 1 1 1 1 1 1 1 ...
 $ CSP.Name : Factor w/ 317 levels "Adur","Allerdale",...: 14 14 14 14 14 14 14 14 14 14 ...
 $ Offence.Description: Factor w/ 131 levels "Absconding from lawful custody",...: 1 2 3 4 5 6 7 8 9 10 ...
 $ Offence.Group : Factor w/ 9 levels "Criminal damage and arson",...: 3 7 7 8 8 8 8 8 3 1 ...
 $ Offence.Subgroup : Factor w/ 25 levels "Arson","Bicycle theft",...: 7 9 9 8 8 5 5 21 7 1 ...
 $ Offence.Code : Factor w/ 131 levels "1", "104", "...: 105 100 102 48 47 41 42 56 103 82 ...
 $ Number.of.offences : int 0 0 0 0 0 0 0 5 2 0 4 ...

> nrow(crimes)
[1] 191260

> ncol(crimes)
[1] 9

> dim(crimes)
[1] 191260 9
```

3. Select 3 columns for the cluster analysis

```
crimes<-crimes[,c(3,7,9)]  
head(crimes)
```

```
> head(crimes)  
      Force.Name      Offence.Subgroup Number.of.offences  
1 Avon & Somerset Miscellaneous crimes                0  
2 Avon & Somerset Other sexual offences                0  
3 Avon & Somerset Other sexual offences                0  
4 Avon & Somerset Non-domestic burglary                0  
5 Avon & Somerset Non-domestic burglary                0  
6 Avon & Somerset      Domestic burglary                0  
> |
```

4. Transform data

reshape2 is an R package written by Hadley Wickham that makes it easy to transform data between wide and long formats.

```
install.packages("reshape2") # install "reshape2" package  
library(reshape2)           # activate "reshape2" package
```

dcast() function can be used to convert from a long format to a wide format.

Usage: The dcast() function takes three arguments:

- data: A molten data frame.
- formula: A formula that specifies how you want to cast the data. This formula takes the form x_variable ~ y_variable.
- fun.aggregate: A function to use if the casting formula results in data aggregation (for example, length(), sum(), or mean()).

```
crimes_pivot <- dcast(crimes, Force.Name ~ Offence.Subgroup, sum,  
                      value.var = "Number.of.offences")  
head(crimes_pivot)
```

	Force.Name	Arson	Bicycle theft	Criminal damage	Death or serious injury - unlawful driving	Domestic burglary	Homicide	Miscellaneous crimes
1	Avon & Somerset	717	3443	15492	3	8029	14	1906
2	Bedfordshire	215	1023	5674	8	4512	4	524
3	Cambridgeshire	355	4295	7818	25	4452	7	998
4	Cheshire	491	1397	9877	8	3489	5	1731
5	Cleveland	368	963	9187	11	4329	3	882
6	Cumbria	193	269	4948	16	1148	3	604

5. Set values in Force.Name column as rownames

```
rownames(crimes_pivot) <- crimes_pivot[,1]
```

```
crimes_pivot[,1] <- NULL
head(crimes_pivot)
```

```
> head(crimes_pivot)
```

	Arson	Bicycle theft	Criminal damage	Death or serious injury - unlawful driving	Domestic burglary	Homicide
Avon & Somerset	717	3443	15492	3	8029	14
Bedfordshire	215	1023	5674	8	4512	4
Cambridgeshire	355	4295	7818	25	4452	7
Cheshire	491	1397	9877	8	3489	5
Cleveland	368	963	9187	11	4329	3
Cumbria	193	269	4948	16	1148	3

6. Normalising the dataset

```
normalise <- function(df)
{
  return(((df- min(df)) / (max(df)-min(df)) * (1-0)) +0)
}

#Normalise the crimes_pivot data set using above normalise function
```

```
Force.Name<-rownames(crimes_pivot)
crimes_pivot_n<-as.data.frame(lapply(crimes_pivot,normalise))
rownames(crimes_pivot_n)<-Force.Name
```

```
head(crimes_pivot_n)
```

```
> head(crimes_pivot_n)
```

	Arson	Bicycle.theft	Criminal.damage	Death.or.serious.injury...unlawful.driving	Domestic.burglary	Homicide	Miscellaneous.crimes
Avon & Somerset	0.27318008	0.162085731	0.26572171	0.03846154	0.14046618	0.077419355	0.18325955
Bedfordshire	0.08084291	0.042022227	0.09454818	0.13461538	0.07882930	0.012903226	0.04097601
Cambridgeshire	0.13448276	0.204356023	0.13192810	0.46153846	0.07777778	0.032258065	0.08977659
Cheshire	0.18659004	0.060577496	0.16782607	0.13461538	0.06090081	0.019354839	0.16524246
Cleveland	0.13946360	0.039045446	0.15579615	0.19230769	0.07562215	0.006451613	0.07783383
Cumbria	0.07241379	0.004614011	0.08189061	0.28846154	0.01987382	0.006451613	0.04921240

7. Assessing clustering tendency

`get_clust_tendency()` function in **factoextra** package helps to calculate **Clustering tendency**. It's determines whether a given dataset contains meaningful clusters (i.e., non-random structure). **Hopkins statistic** is used to assess the **clustering tendency** of a dataset by measuring the probability that a given dataset is generated by a uniform data distribution. In other words, it tests the **spatial randomness** of the data.

If the value of **Hopkins statistic** is close to zero, then we can reject the null hypothesis and conclude that the dataset is significantly a clusterable data.

The null and the alternative hypotheses are defined as follow:

- **Null hypothesis:** the dataset is uniformly distributed (i.e., no meaningful clusters)
- **Alternative hypothesis:** the dataset is not uniformly distributed (i.e., contains meaningful clusters)

```
tendency <- get_clust_tendency(crimes_pivot_n, 30, graph = TRUE)
tendency$hopkins_stat
```


```
> tendency$hopkins_stat
[1] 0.1746332
```

8. If the data is clusterable, then how to choose the right number of expected clusters (k)?

#fviz_nbclust() function determines and visualizes the optimal number of clusters using different methods: within cluster sums of squares, average silhouette and gap statistics.

fviz_nbclust() function plots the Within Cluster Sum of Squares and the number of clusters to find the location of a bend or a knee in the plot which is considered as an indicator of the appropriate number of clusters.

```
fviz_nbclust(crimes_pivot_n, kmeans, method = "wss")
```




9. Perform k-means clustering on a **crimes_pivot_n** dataset with k=3

```
set.seed(123)
km.fit <- kmeans(crimes_pivot_n, 3, nstart = 30)
km.fit$cluster
km.fit$size
```

```
> km.fit$size
[1] 31  1 11
```

10. Visualize clusters using **fviz_cluster()** function in **factoextra** package

```
fviz_cluster(km.fit, crimes_pivot_n)
```



11. Since “Metropolitan Police” has a large number of crimes and behave like an outlier. So, perform k-means clustering on a **crimes_pivot_n** dataset with k=3, but without “Metropolitan Police” crime figures.

```
crimes_pivot_n2 <- subset(crimes_pivot_n,
                           rownames(crimes_pivot_n) != "Metropolitan Police")
```




```
# Compute k-means
set.seed(123)
km.fit2 <- kmeans(crimes_pivot_n2, 3, nstart = 30)
km.fit2$cluster
km.fit2$size
```

```
> km.fit2$size
[1] 26  2 14
```

```
# Visualise the clusters
```

```
fviz_cluster(km.fit2, crimes_pivot_n2)
```




12. Perform k-means clustering on a **crimes_pivot_n** dataset with k=3, but without below mentioned areas.

- Metropolitan Police
- Greater Manchester
- West Midlands
- West Yorkshire

```
crimes_pivot_n3<-subset(crimes_pivot_n, !(rownames(crimes_pivot_n) %in%
c("Metropolitan Police","Greater Manchester",
"West Midlands", "West Yorkshire")))
```

```
# Compute k-means
set.seed(123)
km.fit3 <- kmeans(crimes_pivot_n3, 3, nstart = 30)
km.fit3$cluster
km.fit3$size
```

```
fviz_cluster(km.fit3, crimes_pivot_n3)
```



```
fviz_cluster(km.fit3, crimes_pivot_n3, ellipse.type = "norm")
```

