

ASDM Workshop Week 5: Text Mining with R

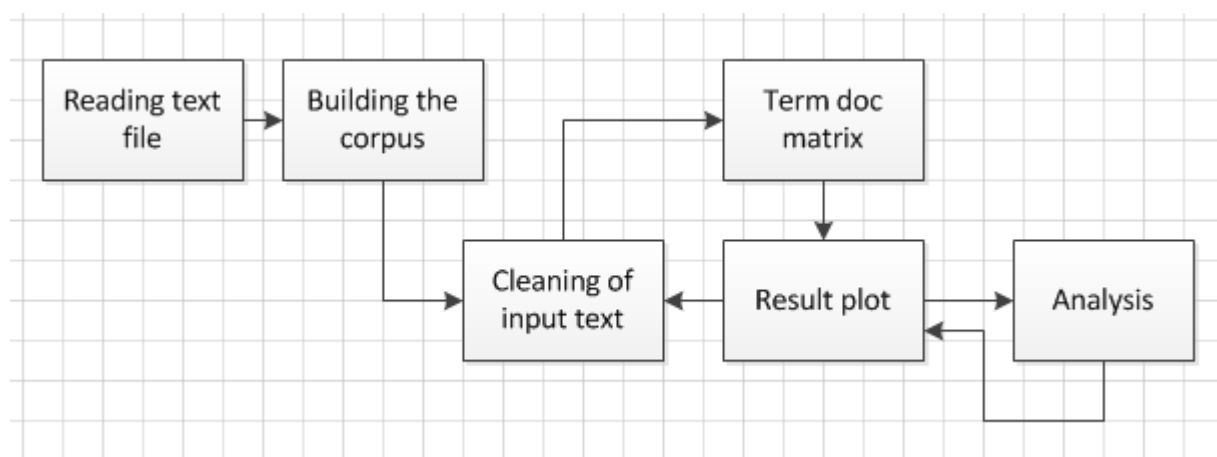
- Charith Silva, Dr. Mo Saraee

Introduction :

Text mining encompasses a vast field of theoretical approaches and methods with one thing in common: text as input information. This allows various definitions, ranging from an extension of classical data mining to texts to more sophisticated formulations like “the use of large online text collections to discover new facts and trends about the world itself” (Hearst 1999). In general, text mining is an interdisciplinary field of activity amongst data mining, linguistics, computational statistics, and computer science. Standard techniques are text classification, text clustering, ontology and taxonomy creation, document summarization and latent corpus analysis. In addition, a lot of techniques from related fields like information retrieval are commonly used (Feinerer et al., 2008).

Text mining process:

Text mining involves a series of activities to be performed in order to efficiently mine the information.



- Reading text file
- Building the corpus
- Cleaning of input text: eg : remove “commas, web links etc”
- Term doc matrix: Since there are so many words in our document how often each word appears we create a matrix for that.
- Result plot: Then we plot the results and go back and forth between analysis, plot and cleaning text.
- Analysis: Analyze the results.

Part 1: Exercise – Text mining with unstructured text data

In this workshop, we will be using “tm” package.

1. Start RStudio.

2. Change working directory.

File → More → Go To Working Directory...

In the Go To Working Directory dialogue, navigate to and select the folder where you saved your data file eg: F:\ASDM\Week5. Click OK.

3. Open a new R script window:

File → New File → R script

4. Install “tm” package in RStudio.

```
#tm package provides a framework for text mining applications
install.packages("tm")
```

5. Activate tm library

```
library(tm)
```

6. Download and Read the data file

We will be using TMwithR.txt data to mine text in the text file. The file can be downloaded from Blackboard.

```
dataset<- readLines("TMwithR.txt")
```

The data is read into the data frame called “dataset”. To see what is in the document use dataset command.

```
dataset
```

You can see there are 424 different items of text.

7. Inspect the dataset in R.

Once the file has been imported to R, we often want to do few things to explore the dataset:

```
names(dataset)
head(dataset)
tail(dataset)
summary(dataset)
str(dataset)
```

8. Use the following code to create corpus.

```
#converting the text file to corpus
#Corpus is collections of documents containing (natural language)
text. Corpus is the main structure that tm uses for storing and
manipulating text documents.
```

```
mycorpus <- Corpus(VectorSource(dataset))
mycorpus
```

Now use the `mycorpus` command to see the documents in corpus.

Result:

```
> mycorpus
<<SimpleCorpus>>
Metadata: corpus specific: 1, document level (indexed): 0
Content: documents: 424
```

9. Use the following commands to inspect the items in corpus.

```
inspect(mycorpus[1])
inspect(mycorpus[2])
inspect(mycorpus[3])
```

```
> inspect(mycorpus[1])
<<SimpleCorpus>>
Metadata: corpus specific: 1, document level (indexed): 0
Content: documents: 1

[1] IEEE TRANSACTIONS ON RELIABILITY, VOL. 58, NO. 4, DECEMBER 2009
> inspect(mycorpus[2])
<<SimpleCorpus>>
Metadata: corpus specific: 1, document level (indexed): 0
Content: documents: 1

[1]
> inspect(mycorpus[3])
<<SimpleCorpus>>
Metadata: corpus specific: 1, document level (indexed): 0
Content: documents: 1

[1] 649
```

```
inspect(mycorpus[8])
mycorpus[8]
```

```
> inspect(mycorpus[8])
<<SimpleCorpus>>
Metadata: corpus specific: 1, document level (indexed): 0
Content: documents: 1

[1] Abstract--Although the need for collecting warranty data originated from financial reasons, it is also extensively used for modeling and analysis to support managerial decision-making in industries. Strategic, tactical, and operational level decisions involving warranty cost very often use warranty spending forecasts that are developed using statistical methods. Existing literature provides warranty forecasting approaches involving variables such as mileage accumulation rate, failure rate, repeat repair rate, and cost per repair. However, there are several key failure modes that are known to be influenced by seasonality. For example, 'engine slow to start' conditions drive a higher claim rate in colder months than in warmer months. Accommodation of such failure modes influenced by seasonality has not been considered in the warranty cost modeling literature. This paper presents a flexible approach for developing a monthly warranty spend forecasting model that incorporates calendar month seasonality, business days per month for authorized service centers, and sales ramp-up in addition to the earlier mentioned variables. On one hand, the model allows development of warranty spend forecasts for entire warranty coverage to support strategic level decisions; on the other hand, forecasts for monthly warranty spend help support tactical and operational level decisions. The workability of the proposed methodology is illustrated using an application example.
> mycorpus[8]
<<SimpleCorpus>>
Metadata: corpus specific: 1, document level (indexed): 0
Content: documents: 1
```

10. Use the following command to clean up and convert the document in lower alphabets.

#tm_map function can be used as interface to apply transformation functions to corpus.

```
mycorpus <- tm_map(mycorpus,tolower)
inspect(mycorpus[8])
```

Result:

The document has been converted to lowercase.

```
> mycorpus = tm_map(mycorpus,tolower)
> inspect(mycorpus[8])
<<SimpleCorpus>>
Metadata: corpus specific: 1, document level (indexed): 0
Content: documents: 1

[1] abstract--although the need for collecting warranty data originated from financial reasons, it is also extensively used for modeling and analysis to support managerial decision-making in industries. strategic, tactical, and operational level decisions involving warranty cost very often use warranty spending forecasts that are developed using statistical methods. existing literature provides warranty forecasting approaches involving variables such as mileage accumulation rate, failure rate, repeat repair rate, and cost per repair. however, there are several key failure modes that are known to be influenced by seasonality. for example, 'engine slow to start' conditions drive a higher claim rate in colder months than in warmer months. accommodation of such failure modes influenced by seasonality has not been considered in the warranty cost modeling literature. this paper presents a flexible approach for developing a monthly warranty spend forecasting model that incorporates calendar month seasonality, business days per month for authorized service centers, and sales ramp-up in addition to the earlier mentioned variables. on one hand, the model allows development of warranty spend forecasts for entire warranty coverage to support strategic level decisions; on the other hand, forecasts for monthly warranty spend help support tactical and operational level decisions. the workability of the proposed methodology is illustrated using an application example.
```

11. Use `getTransformations()` function to retrieve the list of predefined transformations (mappings) which can be used with `tm_map` function.

```
getTransformations()
```

```
> getTransformations()
[1] "removeNumbers" "removePunctuation" "removeWords" "stemDocument" "stripWhitespace"
```

12. Use the following command to remove punctuations.

```
mycorpus <-tm_map(mycorpus,removePunctuation)
inspect(mycorpus[8])
```

Result:

```
> mycorpus=tm_map(mycorpus,removePunctuation)
> inspect(mycorpus[8])
<<SimpleCorpus>>
Metadata: corpus specific: 1, document level (indexed): 0
Content: documents: 1

[1] abstractalthough the need for collecting warranty data originated from financial reason
s it is also extensively used for modeling and analysis to support managerial decisionmakin
g in industries strategic tactical and operational level decisions involving warranty cost
very often use warranty spending forecasts that are developed using statistical methods exi
sting literature provides warranty forecasting approaches involving variables such as milea
ge accumulation rate failure rate repeat repair rate and cost per repair however there are
several key failure modes that are known to be influenced by seasonality for example engine
slow to start conditions drive a higher claim rate in colder months than in warmer months
accommodation of such failure modes influenced by seasonality has not been considered in th
e warranty cost modeling literature this paper presents a flexible approach for developing
a monthly warranty spend forecasting model that incorporates calendar month seasonality bus
iness days per month for authorized service centers and sales rampup in addition to the ear
lier mentioned variables on one hand the model allows development of warranty spend forecas
ts for entire warranty coverage to support strategic level decisions on the other hand fore
casts for monthly warranty spend help support tactical and operational level decisions the
workability of the proposed methodology is illustrated using an application example
```

Similarly we can remove numbers using following line of code.

```
mycorpus <-tm_map(mycorpus,removeNumbers)
```

13. We can also remove stop words from the document.

Use the following line of code to see the stop words in English language.

```
stopwords("en")
```

```
> stopwords("en")
[1] "i" "me" "my" "myself" "we" "our"
[2] "ours" "ourselves" "you" "your" "yours" "yourself"
[3] "yourselves" "he" "him" "his" "himself" "she"
[4] "her" "hers" "herself" "it" "its" "itself"
[5] "they" "them" "their" "theirs" "themselves" "what"
[6] "which" "who" "whom" "this" "that" "these"
[7] "those" "am" "is" "are" "was" "were"
[8] "be" "been" "being" "have" "has" "had"
[9] "having" "do" "does" "did" "doing" "would"
[10] "should" "could" "ought" "i'm" "you're" "he's"
[11] "she's" "it's" "we're" "they're" "i've" "you've"
[12] "we've" "they've" "i'd" "you'd" "he'd" "she'd"
[13] "we'd" "they'd" "i'll" "you'll" "he'll" "she'll"
[14] "we'll" "they'll" "isn't" "aren't" "wasn't" "weren't"
[15] "hasn't" "haven't" "hadn't" "doesn't" "don't" "didn't"
[16] "won't" "wouldn't" "shan't" "shouldn't" "can't" "cannot"
[17] "couldn't" "mustn't" "let's" "that's" "who's" "what's"
[18] "here's" "there's" "when's" "where's" "why's" "how's"
[19] "a" "an" "the" "and" "but" "if"
[20] "or" "because" "as" "until" "while" "of"
[21] "at" "by" "for" "with" "about" "against"
[22] "between" "into" "through" "during" "before" "after"
[23] "above" "below" "to" "from" "up" "down"
[24] "in" "on" "off" "over" "under"
[25] "again" "out" "once" "then" "here" "there"
[26] "when" "where" "why" "how" "all" "any"
[27] "both" "each" "few" "more" "most" "other"
[28] "some" "such" "no" "nor" "not" "only"
```

Use the following line of code to remove stop words:

```
dataclean <-tm_map(mycorpus,removeWords,stopwords("english"))
inspect(dataclean[8])
```

Result:

```
> dataclean <-tm_map(mycorpus,removeWords,stopwords("english"))
> inspect(dataclean[8])
<<SimpleCorpus>>
Metadata: corpus specific: 1, document level (indexed): 0
Content: documents: 1

[1] abstractalthough need collecting warranty data originated financial reasons also e
xtensively used modeling analysis support managerial decisionmaking industries strategi
c tactical operational level decisions involving warranty cost often use warranty spendin
g forecasts developed using statistical methods existing literature provides warranty for
ecasting approaches involving variables mileage accumulation rate failure rate repeat rep
air rate cost per repair however several key failure modes known influenced seasona
lity example engine slow start conditions drive higher claim rate colder months warme
r months accommodation failure modes influenced seasonality considered warranty cos
t modeling literature paper presents flexible approach developing monthly warranty spen
d forecasting model incorporates calendar month seasonality business days per month autho
rized service centers sales rampup addition earlier mentioned variables one hand mode
l allows development warranty spend forecasts entire warranty coverage support strategic
level decisions hand forecasts monthly warranty spend help support tactical operation
al level decisions workability proposed methodology illustrated using application exam
ple
```

14. As you can see there are lots of white spaces in the document. Use the following line of code to remove those white spaces.

```
dataclean <- tm_map(dataclean,stripWhitespace)
inspect(dataclean[8])
```

Result:

```
> dataclean <- tm_map(dataclean,stripWhitespace)
> inspect(dataclean[8])
<<SimpleCorpus>>
Metadata: corpus specific: 1, document level (indexed): 0
Content: documents: 1

[1] abstractalthough need collecting warranty data originated financial reasons also extens
ively used modeling analysis support managerial decisionmaking industries strategic tactica
l operational level decisions involving warranty cost often use warranty spending forecasts
developed using statistical methods existing literature provides warranty forecasting appr
oaches involving variables mileage accumulation rate failure rate repeat repair rate cost p
er repair however several key failure modes known influenced seasonality example engine slo
w start conditions drive higher claim rate colder months warmer months accommodation failur
e modes influenced seasonality considered warranty cost modeling literature paper presents
flexible approach developing monthly warranty spend forecasting model incorporates calendar
month seasonality business days per month authorized service centers sales rampup addition
earlier mentioned variables one hand model allows development warranty spend forecasts ent
ire warranty coverage support strategic level decisions hand forecasts monthly warranty spe
nd help support tactical operational level decisions workability proposed methodology illus
trated using application example
```

15. The next step is to create a Document-Term Matrix (DTM). DTM is a matrix that lists all occurrences of words in the corpus. In DTM, documents are represented by rows and the terms (or words) by columns. If a word occurs in a particular document n times, then the matrix entry for corresponding to that row and column is n , if it doesn't occur at all, the entry is 0.

Use the following line of code to create the term document matrix

```
dtm <- TermDocumentMatrix(dataclean,
                           control = list(minWordLength=c(1,Inf))
                           )
dtm
```

```
> dtm
<<TermDocumentMatrix (terms: 718, documents: 424)>>
Non-/sparse entries: 2406/302026
Sparsity           : 99%
Maximal term length: 22
Weighting          : term frequency (tf)
```

#findFreqTerms function can be used to find frequent terms in a document-term or term-document matrix.

```
findFreqTerms(dtm, lowfreq = 2)
```

Result:

```
> dtm=TermDocumentMatrix(dataclean,control = list(minWordLength=c(1,Inf)))
> findFreqTerms(dtm,lowfreq = 2)
[1] "2009"          "december"      "ieeee"         "reliability"
[5] "transactions"  "vol"          "calendar"      "failures"
[9] "forecasting"   "influenced"    "month"         "seasonality"
[13] "spend"         "subsystem"     "warranty"      "rai"
[17] "accumulation"  "addition"      "also"          "analysis"
[21] "application"   "approach"      "approaches"    "authorized"
[25] "business"      "centers"       "claim"         "conditions"
[29] "cost"          "coverage"      "data"          "days"
[33] "decisions"     "developed"     "developing"    "engine"
[37] "example"       "extensively"   "failure"       "flexible"
[41] "forecasts"     "hand"         "help"          "higher"
[45] "however"       "illustrated"   "incorporates"  "involving"
[49] "key"           "known"        "level"         "literature"
[53] "methodology"   "methods"      "mileage"       "model"
[57] "modeling"      "modes"        "monthly"       "months"
[61] "need"          "often"        "one"           "operational"
[65] "paper"         "per"          "presents"      "proposed"
[69] "provides"      "rampup"       "rate"          "repair"
```

16. Use the following line of code to see the words with frequency in the document matrix.

```
termFrequency <- rowSums(as.matrix(dtm))
termFrequency
```

Result:

```
> termFrequency = rowSums(as.matrix(dtm))
> termFrequency
```

2009	16	december	15	ieee	14	reliability	8	transactions	5
vol	5		649	calendar	7	failures	5	forecasting	21
influenced	15	month	75	seasonality	32	spend	9	subsystem	15
warranty	120	bharatendra	1	rai	8	abstractalthough	1	accommodation	1
accumulation	15	addition	2	allows	1	also	9	analysis	5
application	4	approach	4	approaches	3	authorized	12	business	35
centers	2	claim	33	colder	1	collecting	1	conditions	3
considered	1	cost	46	coverage	8	data	18	days	34
decisionmaking	1	decisions	18	developed	8	developing	7	development	1

Use the following line of code to see the words with frequency greater than 15.


```
termFrequency <- subset(termFrequency,termFrequency>=15)
termFrequency
```

```
> termFrequency=subset(termFrequency,termFrequency>=15)
> termFrequency
```

2009	december	forecasting	influenced	month	seasonality	subsystem	warranty
16	15	21	15	75	32	15	120
accumulation	business	claim	cost	data	days	decisions	failure
15	35	33	46	18	34	18	34
level	mileage	model	monthly	months	per	rampup	rate
19	27	56	23	20	35	17	35
repair	sales	used	using	year	first	claims	number
25	18	15	22	44	22	55	31
vehicles	vehicle	production	fig	can	estimates	rates	baseline
43	30	32	25	21	32	15	18
indices	obtained	seasonal					
22	15	16					

17. Plot results with barplot() function

```
barplot(termFrequency,las=2,col=rainbow(20))
```



18. Finally to see the word cloud of the words, use the “wordcloud” package.

URL: <https://cran.r-project.org/web/packages/wordcloud/wordcloud.pdf>

```
install.packages("wordcloud")    #install "wordcloud"
library(wordcloud)                #load "wordcloud"
```

Use following lines of code to see the word cloud.

```
wordfreq <- sort(termFrequency, decreasing = TRUE)
wordfreq
```

```
> wordfreq<-sort(termFrequency,decreasing = TRUE)
> wordfreq
      warranty      month      model      claims
      120         89         56         55
      cost         year      vehicles      business
      46          44         43         35
      per          rate      days         failure
      35          35         34         34
      claim      seasonality      production      estimates
      33          32         32         32
      number      vehicle      mileage      repair
      31          30         27         25
      fig         monthly      using         first
      25          23         22         22
      indices      forecasting      can         months
      22          21         21         20
```

Description of the `wordcloud()` function

Description

Plot a word cloud

Usage

```
wordcloud(words,freq,scale=c(4,.5),min.freq=3,max.words=Inf,
random.order=TRUE, random.color=FALSE, rot.per=.1,
colors="black",ordered.colors=FALSE,use.r.layout=FALSE,
fixed.asp=TRUE, ...)
```

Arguments

words	the words
freq	their frequencies
scale	A vector of length 2 indicating the range of the size of the words.
min.freq	words with frequency below min.freq will not be plotted
max.words	Maximum number of words to be plotted. least frequent terms dropped
random.order	plot words in random order. If false, they will be plotted in decreasing frequency
random.color	choose colors randomly from the colors. If false, the color is chosen based on the frequency
rot.per	proportion words with 90 degree rotation
colors	color words from least to most frequent
ordered.colors	if true, then colors are assigned to words in order
use.r.layout	if false, then c++ code is used for collision detection, otherwise R is used
fixed.asp	if TRUE, the aspect ratio is fixed. Variable aspect ratio only supported if rot.per==0
...	Additional parameters to be passed to text (and strheight,strwidth).

Details


If freq is missing, then words can either be a character vector, or Corpus. If it is a vector and freq is missing, standard stop words will be removed prior to plotting.

```
wordcloud(words = names(wordfreq),
          freq=wordfreq,max.words=100,
          min.freq = 5,
          random.order = F)
```



There are many different variations of word cloud.

```
wordcloud(words = names(wordfreq),
          freq=wordfreq,max.words=100,
          min.freq = 5,
          random.order = F,
          colors = rainbow(20))
```



[illegible][illegible]

[illegible]

```
#We now see the distribution of the 50 most frequent words in a
barplot.
```

```
barplot(wordfreq[1:50],
       xlab = "term",
       ylab = "frequency",
       las=2,
       col=heat.colors(50))
```



Part 2: Exercise - Text mining and clustering using Twitter data

In this tutorial, we will use Twitter data to compare the trend of popular health related news from 3 news agencies.

Twitter is a popular micro blogging service where users create status messages (called "tweets"). Tweets are short messages with a maximum length of 140 characters. The distinguishing characteristics of tweets are hashtags. Hashtags are used for logically grouping tweets and searching them. Authors of those messages write about their life, share opinions on variety of topics and discuss current issues. As more and more users post about products and services they use, or express their political, religious views, etc... Micro blogging websites become valuable sources of people's opinions and sentiments. Such data can be efficiently used for marketing or social studies. As a result, there has been a tremendous need to design methods and algorithms which can effectively process a wide variety of text applications (source : Suprajha S,Yogitha C :A Study on Sentiment Analysis using Tweeter Data).

Dataset Information:

The data was collected in 2015 using Twitter API. Each file is related to one Twitter account of a news agency. For example, bbchealth.csv is related to BBC health news. Each line of the dataset contains tweet id, date and time and tweet. This text data has been used to evaluate the performance of topic models on short text data. However, it can be used for other tasks such as clustering.

```
#Read the data file
bbchealth <- read.csv("bbchealth.csv", header= TRUE)
cnnhealth <- read.csv("cnnhealth.csv", header= TRUE)
foxhealth <- read.csv("foxnewshealth.csv", header= TRUE)
```

```
#Inspect the dataset
head(bbchealth)
head(cnnhealth)
head(foxhealth)
```

```
> head(bbchealth)
  twee_id      date.and.time      tweet
1 5.86e+17 Thu Apr 09 01:31:50 +0000 2015 Breast cancer risk test devised http://bbc.in/1CimpJF
2 5.86e+17 Wed Apr 08 23:30:18 +0000 2015 GP workload harming care - BMA poll http://bbc.in/1ChTBRv
3 5.86e+17 Wed Apr 08 23:30:18 +0000 2015 Short people's 'heart risk greater' http://bbc.in/1ChTANp
4 5.86e+17 Wed Apr 08 18:05:28 +0000 2015 New approach against HIV 'promising' http://bbc.in/1E6jAjt
5 5.86e+17 Wed Apr 08 13:19:33 +0000 2015 Coalition 'undermined NHS' - doctors http://bbc.in/1CnLwK7
6 5.86e+17 Wed Apr 08 09:18:39 +0000 2015 Review of case against NHS manager http://bbc.in/1Ffj6ci
```

```
#Inspect the tweet column in the datasets
head(bbchealth$tweet)
head(cnnhealth$tweet)
```

```
head(foxhealth$tweet)
```

```
#create text vectors
```

```
bbchealth_tweet<- bbchealth$tweet
```

```
cnnhealth_tweet<- cnnhealth$tweet
```

```
foxhealth_tweet<- foxhealth$tweet
```

```
#convert all text to lower case
```

```
bbchealth_tweet<- tolower(bbchealth_tweet)
```

```
cnnhealth_tweet<- tolower(cnnhealth_tweet)
```

```
foxhealth_tweet<- tolower(foxhealth_tweet)
```

`gsub()` function replaces all matches of a string, if the parameter is a string vector, returns a string vector of the same length and with the same attributes (after possible coercion to character). Elements of string vectors which are not substituted will be returned unchanged (including any declared encoding). `gsub()` function can use regular expressions as search string.

What is a regular expression?

Regular expression is a pattern that describes a set of strings. Simply speaking, regular expression is an "instruction" given to a function on what and how to match or replace strings.

Additional reading materials about regular expression

- https://rstudio-pubs-static.s3.amazonaws.com/74603_76cd14d5983f47408fdf0b323550b846.html
- <http://biostat.mc.vanderbilt.edu/wiki/pub/Main/SvetlanaEdenRFiles/regExprTalk.pdf>

Some popular regular expression syntaxs:

Syntax	Description
<code>\\d</code>	Digit, 0,1,2 ... 9
<code>\\D</code>	Not Digit
<code>\\s</code>	Space
<code>\\S</code>	Not Space
<code>\\w</code>	Word
<code>\\W</code>	Not Word
<code>\\t</code>	Tab
<code>\\n</code>	New line
<code>^</code>	Beginning of the string
<code>\$</code>	End of the string
<code>\\</code>	Escape special characters, e.g. <code>\\</code> is "\", <code>\\+</code> is "+"
<code> </code>	Alternation match. e.g. <code>/(e d)n/</code> matches "en" and "dn"
<code>.</code>	Any character, except <code>\\n</code> or line terminator

[ab]	a or b
[^ab]	Any character except a and b
[0-9]	All Digit
[A-Z]	All uppercase A to Z letters
[a-z]	All lowercase a to z letters
[A-z]	All Uppercase and lowercase a to z letters
i+	i at least one time
i*	i zero or more times
i?	i zero or 1 time
i{n}	i occurs n times in sequence
i{n1,n2}	i occurs n1 - n2 times in sequence
i{n1,n2}?	non greedy match, see above example
i{n,}	i occurs >= n times
[:alnum:]	Alphanumeric characters: [:alpha:] and [:digit:]
[:alpha:]	Alphabetic characters: [:lower:] and [:upper:]
[:blank:]	Blank characters: e.g. space, tab
[:cntrl:]	Control characters
[:digit:]	Digits: 0 1 2 3 4 5 6 7 8 9
[:graph:]	Graphical characters: [:alnum:] and [:punct:]
[:lower:]	Lower-case letters in the current locale
[:print:]	Printable characters: [:alnum:], [:punct:] and space
[:punct:]	Punctuation character: ! " # \$ % & ' () * + , - . / : ; < = > ? @ [\] ^ _ ` { } ~
[:space:]	Space characters: tab, newline, vertical tab, form feed, carriage return, space
[:upper:]	Upper-case letters in the current locale
[:xdigit:]	Hexadecimal digits: 0 1 2 3 4 5 6 7 8 9 A B C D E F a b c d e f

```
#Replace blank space ("rt")
bbchealth_tweet <- gsub("rt", "", bbchealth_tweet)
cnnhealth_tweet <- gsub("rt", "", cnnhealth_tweet)
foxhealth_tweet <- gsub("rt", "", foxhealth_tweet)

#Replace tweeter @UserName
bbchealth_tweet <- gsub("@\\w+", "", bbchealth_tweet)
cnnhealth_tweet <- gsub("@\\w+", "", cnnhealth_tweet)
foxhealth_tweet <- gsub("@\\w+", "", foxhealth_tweet)

#Replace links in the tweets
bbchealth_tweet <- gsub("http\\S+\\s*", "", bbchealth_tweet)
cnnhealth_tweet <- gsub("http\\S+\\s*", "", cnnhealth_tweet)
foxhealth_tweet <- gsub("http\\S+\\s*", "", foxhealth_tweet)

#Remove punctuation
bbchealth_tweet <- gsub("[[:punct:]]", "", bbchealth_tweet)
cnnhealth_tweet <- gsub("[[:punct:]]", "", cnnhealth_tweet)
foxhealth_tweet <- gsub("[[:punct:]]", "", foxhealth_tweet)
```

```
#Remove tabs
bbchealth_tweet <- gsub("[ |\t]{2,}", "", bbchealth_tweet)
cnnhealth_tweet <- gsub("[ |\t]{2,}", "", cnnhealth_tweet)
foxhealth_tweet <- gsub("[ |\t]{2,}", "", foxhealth_tweet)
```

```
#Remove "video" word in the tweets
bbchealth_tweet <- gsub("video", "", bbchealth_tweet)
cnnhealth_tweet <- gsub("video", "", cnnhealth_tweet)
foxhealth_tweet <- gsub("video", "", foxhealth_tweet)
```

```
#Remove blank spaces at the beginning
bbchealth_tweet <- gsub("^ ", "", bbchealth_tweet)
cnnhealth_tweet <- gsub("^ ", "", cnnhealth_tweet)
foxhealth_tweet <- gsub("^ ", "", foxhealth_tweet)
```

```
#Remove blank spaces at the end
bbchealth_tweet <- gsub(" $", "", bbchealth_tweet)
cnnhealth_tweet <- gsub(" $", "", cnnhealth_tweet)
foxhealth_tweet <- gsub(" $", "", foxhealth_tweet)
```

```
#Inspect the vectors after cleaning
head(bbchealth_tweet)
head(cnnhealth_tweet)
head(foxhealth_tweet)
```

```
> head(bbchealth_tweet)
[1] "breast cancer risk test devised" "gp workload harming carebma poll" "sho peoples hea risk greater"
[4] "new approach against hiv promising" "coalition undermined nhsdoctors" "review of case against nhs manager"
> head(cnnhealth_tweet)
[1] "an abundance of online info can turn us into ehypochondriacs or worse lead us to neglect getting the care we need"
[2] "a plantbased diet that incorporates fish may be the key to preventing colorectal cancers"
[3] "it doesnt take much to damage your hearing at a spos bar or nightclub thats why a billion people are at risk"
[4] "forever young discover this islandâ€™s secrets to longevity on thewonderlist w"
[5] "is posttraumatic stress disorder in your genes a simple blood test may one day help tell you"
[6] "maysoon zayid a touring standup comic with cerebral palsy has a message to share"
> head(foxhealth_tweet)
[1] "injury prevention programs unpopular with high school coaches" "6 dietary changes to make midlife"
[3] "massachusetts governor gets head shaved to suppo charity" "dad wins 3 marathons in 8 days winnings to help ailing son"
[5] "possible cure for melanoma" "wear orange glasses to get better sleep study says"
```

```
library(tm) #load tm package
```

```
#converting the text vectors to corpus
bbchealth_corpus <- Corpus(VectorSource(bbchealth_tweet))
bbchealth_corpus
```



```
cnnhealth_corpus <- Corpus(VectorSource(cnnhealth_tweet))
cnnhealth_corpus
```

```
foxhealth_corpus <- Corpus(VectorSource(foxhealth_tweet))
foxhealth_corpus
```

```
> bbchealth_corpus
<<SimpleCorpus>>
Metadata: corpus specific: 1, document level (indexed): 0
Content: documents: 2000
> cnnhealth_corpus <- Corpus(VectorSource(cnnhealth_tweet))
> cnnhealth_corpus
<<SimpleCorpus>>
Metadata: corpus specific: 1, document level (indexed): 0
Content: documents: 2000
> foxhealth_corpus <- Corpus(VectorSource(foxhealth_tweet))
> foxhealth_corpus
<<SimpleCorpus>>
Metadata: corpus specific: 1, document level (indexed): 0
Content: documents: 2000
```

```
#clean up corpus by removing stop words, number and Whitespace
```

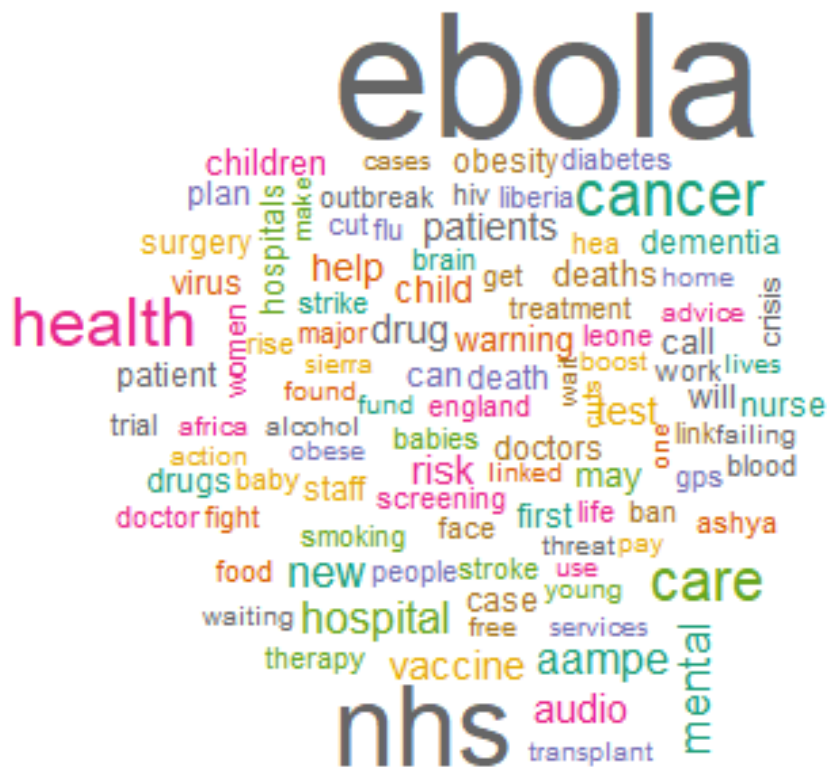
```
bbchealth_corpus <- tm_map(bbchealth_corpus,
                           removeWords, stopwords("english"))
bbchealth_corpus <- tm_map(bbchealth_corpus, removeNumbers)
bbchealth_corpus <- tm_map(bbchealth_corpus, stripWhitespace)
inspect(bbchealth_corpus )
```

```
cnnhealth_corpus <- tm_map(cnnhealth_corpus,
                           removeWords, stopwords("english"))
cnnhealth_corpus <- tm_map(cnnhealth_corpus, removeNumbers)
cnnhealth_corpus <- tm_map(cnnhealth_corpus, stripWhitespace)
inspect(cnnhealth_corpus )
```


```
foxhealth_corpus <- tm_map(foxhealth_corpus ,
                           removeWords, stopwords("english"))
foxhealth_corpus <- tm_map(foxhealth_corpus , removeNumbers)
foxhealth_corpus <- tm_map(foxhealth_corpus , stripWhitespace)
inspect(foxhealth_corpus )
```

```
library("wordcloud") #load wordcloud package


#generate wordclouds
wordcloud(bbchealth_corpus,
  min.freq = 10,
  colors=brewer.pal(8, "Dark2"),
  random.color = TRUE,
  max.words = 100)
```



```
wordcloud(cnnhealth_corpus,
  min.freq = 10,
  colors=brewer.pal(8, "Dark2"),
  random.color = TRUE,
  max.words = 100)
```



```
wordcloud(foxhealth_corpus,
  min.freq = 10,
  colors=brewer.pal(8, "Dark2"),
  random.color = TRUE,
  max.words = 100)
```



```
#Create document-term matrix
```

```
bbchealth_dtm <- DocumentTermMatrix(  
  bbchealth_corpus,  
  control = list(minWordLength=c(3,Inf),  
                 bounds = list(global = c(40, Inf)))  
)
```

```
cnnhealth_dtm <- DocumentTermMatrix(  
  cnnhealth_corpus,  
  control = list(minWordLength=c(3,Inf),  
                 bounds = list(global = c(40, Inf)))  
)
```

```
foxhealth_dtm <- DocumentTermMatrix(  
  foxhealth_corpus,  
  control = list(minWordLength=c(3,Inf),  
                 bounds = list(global = c(40, Inf)))  
)
```

```
bbchealth_dtm  
cnnhealth_dtm  
foxhealth_dtm
```

```
> bbchealth_dtm  
<<DocumentTermMatrix (documents: 2000, terms: 11)>>  
Non-/sparse entries: 1109/20891  
Sparsity           : 95%  
Maximal term length: 8  
Weighting          : term frequency (tf)  
> cnnhealth_dtm  
<<DocumentTermMatrix (documents: 2000, terms: 33)>>  
Non-/sparse entries: 2095/63905  
Sparsity           : 97%  
Maximal term length: 10  
Weighting          : term frequency (tf)  
> foxhealth_dtm  
<<DocumentTermMatrix (documents: 2000, terms: 29)>>  
Non-/sparse entries: 2200/55800  
Sparsity           : 96%  
Maximal term length: 8  
Weighting          : term frequency (tf)
```

```
bbchealth_dtm2 <- as.matrix(bbchealth_dtm)  
cnnhealth_dtm2 <- as.matrix(cnnhealth_dtm)  
foxhealth_dtm2 <- as.matrix(foxhealth_dtm)
```

```
#K-means clustering
```

```
library(cluster) #load cluster package  
library(factoextra) #load factoextra package
```

```

head(bbchealth_dtm2)
bbc_dist <- dist(t(bbchealth_dtm2), method="euclidian")
kfit <- kmeans(bbc_dist, 3)
bbc_dist
kfit
fviz_cluster(kfit,bbc_dist)

```

```

> head(bbchealth_dtm2)
  Terms
Docs cancer new nhs care aampe health hospital drug mental ebola audio
1      1    0  0  0    0    0    0    0    0    0    0    0
2      2    0  0  0  0    0    0    0    0    0    0    0
3      3    0  0  0  0    0    0    0    0    0    0    0
4      4    0  1  0  0    0    0    0    0    0    0    0
5      5    0  0  0  0    0    0    0    0    0    0    0
6      6    0  0  1  0    0    0    0    0    0    0    0
> bbc_dist <- dist(t(bbchealth_dtm2), method="euclidian")
> kfit <- kmeans(bbc_dist, 3)
> bbc_dist
      cancer      new      nhs      care      aampe      health hospital      drug      mental      ebola
new      11.874342
nhs      16.401219 15.620499
care     12.961481 11.704700 16.643317
aampe    11.789826 10.000000 15.748016 11.269428
health   13.928388 12.369317 17.349352 13.190906 12.288206
hospital 11.958261 10.392305 16.000000 11.180340 10.099505 12.529964
drug     10.488088  9.539392 15.198684 11.135529  9.539392 12.000000  9.949874
mental   12.124356 10.295630 16.062378 11.180340 10.198039  7.280110 10.488088  9.949874
ebola    20.591260 19.364917 22.825424 20.248457 19.519221 20.639767 19.621417 18.973666 19.723083
audio    11.401754  9.643651 15.588457 10.954451  9.219544 12.000000  9.643651  9.055385  9.746794 19.235384

```

```

> kfit
K-means clustering with 3 clusters of sizes 1, 1, 9

Cluster means:
      cancer      new      nhs      care      aampe      health hospital      drug      mental      ebola      audio
1 20.59126 19.364917 22.82542 20.24846 19.519221 20.63977 19.621417 18.973666 19.723083 0.000000 19.235384
2 16.40122 15.620499  0.00000 16.64332 15.748016 17.34935 16.000000 15.198684 16.062378 22.82542 15.588457
3 10.72517  9.535482 16.06799 10.39746  9.378216 10.62077  9.582443  9.073059  9.029248 19.76857  9.073915


Clustering vector:
      cancer      new      nhs      care      aampe      health hospital      drug      mental      ebola      audio
3          3          3          2          3          3          3          3          3          1          3

Within cluster sum of squares by cluster:
[1] 0.000 0.000 1058.825
(between_SS / total_SS = 62.0 %)

Available components:

[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss" "betweenss"    "size"
[8] "iter"         "ifault"


```



```

head(cnnhealth_dtm2)
cnn_dist <- dist(t(cnnhealth_dtm2), method="euclidian")
kfit <- kmeans(cnn_dist, 3)
cnn_dist
kfit
fviz_cluster(kfit,cnn_dist)

```



```

head(foxhealth_dtm2)
fox_dist <- dist(t(foxhealth_dtm2), method="euclidian")
kfit <- kmeans(fox_dist, 3)
fox_dist
kfit
fviz_cluster(kfit,fox_dist)

```

